

# VAEPP: Variational Autoencoder with a Pull-back Prior

## Abstract

Many approaches to training generative models by distinct training objectives are proposed in the past. Variational Autoencoder (VAE) is an outstanding model of them based on log-likelihood. In this paper, we propose a novel Pull-back Prior by adjusting the density of prior by a discriminator that can assess the quality of data for VAE. It involves the discriminator from theory of GAN to enrich prior. Based on it, we propose a more general framework, VAE with Pull-back Prior (VAEPP), which uses the existing techniques of VAE and WGAN to improve the log-likelihood, sampling quality and stability of training. VAEPP reaches outstanding NLL and comparable FID on MNIST, Static-MNIST, Fashion-MNIST, Omniglot, CIFAR-10 and CelebA.

## 1 Introduction

How to learn deep generative models that are able to capture complex data pattern in high dimension space, *e.g.*, image datasets, is one of the major challenges in machine learning. Many approaches to training generative models by distinct training objectives are proposed in the past, *e.g.*, Generative Adversarial Networks (GAN) [Goodfellow *et al.*, 2014], Flow-based models [Dinh *et al.*, 2016; Kingma and Dhariwal, 2018], PixelCNN [Van den Oord *et al.*, 2016], and Variational Autoencoders (VAE) [Kingma and Welling, 2014; Rezende *et al.*, 2014].

VAE uses the variational inference and re-parameterization trick to optimize the evidence lower bound of log-likelihood (ELBO). In the past, many researches focused on enriching variational posterior [Kingma *et al.*, 2016; Tomczak and Welling, 2016], but recently some researches show that the standard Gaussian prior could lead to underfitting in latent space, harmful to the performance of VAE [Tomczak and Welling, 2018]. To enrich prior, several learnable priors are proposed [Tomczak and Welling, 2018; Bauer and Mnih, 2019; Takahashi *et al.*, 2019]. Most of them focus on the approximating aggregated posterior which is the integral of variational posterior and is shown as the optimal prior to minimize ELBO. However, existing methods based on aggregated posterior reaches limited performance and the practical

meaning of aggregated posterior is blurry. Recently, we notice that the discriminator can assess the quality of data and **we argue that it is advisable to adjust learnable prior by a discriminator which has more clear practical meaning than approximation for aggregated posterior.**

We propose Pull-back Prior, based on discriminator and learnable prior. Firstly, a discriminator  $D(x)$  is trained for assessing the quality of images. Then, this discriminator is *pulled back* to latent space, defined by  $D(G(z))$ , where  $G(z)$  is the generator. Finally, we adjust prior distribution according to the pull-back discriminator.

We propose a training algorithm for VAE with Pull-back Prior (VAEPP), based on SGVB [Kingma and Welling, 2014] and WGAN gradient penalty terms, which mixes discriminator and gradient penalty term into VAE. We extend it to a more general VAE framework. We also use Langevin Dynamics to improve the sampling of quality. Thanks to the gradient penalty term of WGAN-GP [Gulrajani *et al.*, 2017] and WGAN-div [Wu *et al.*, 2018], and the practical implementation of Langevin dynamics in MEG [Kumar *et al.*, 2019], we enjoy stable efficient training and sampling process.

The main contributions of this paper are the following:

- We propose novel and powerful Pull-back Prior, which is adjusted by a discriminator that can assess the quality of data.
- We propose VAEPP framework to use existing techniques of VAE, *e.g.*, flow posterior, and WGAN, *e.g.*, gradient penalty strategy, to improve the log-likelihood, sampling quality and stability of training.
- In log-likelihood metrics, VAEPP outperforms the models without autoregressive components and is comparable to the autoregressive models on vast common datasets. In FID and IS metrics, it is comparable to GANs and SOTA of VAE on vast common datasets.

## 2 Background

### 2.1 VAE and learnable prior

Many generative models aim to minimize the KL-divergence between empirical distribution  $p^*(x)$  and model distribution  $p_\theta(x)$ , which leads to maximization of log-likelihood. VAE [Kingma and Welling, 2014] models the joint distribution  $p_\theta(x, z)$  and the marginal distribution is  $p_\theta(x) =$

$\int p_\theta(x, z) dz$ . VAE apply variational inference to obtain the evidence lower bound objective (ELBO):

$$\ln p_\theta(x) \geq \mathbb{E}_{q_\phi(z|x)} [\ln p_\theta(x|z) + \ln p_\theta(z) - \ln q_\phi(z|x)] \triangleq \mathcal{L}(x; \theta, \phi) \quad (1)$$

where  $q_\phi(z|x)$  is variational posterior (encoder) and  $p_\theta(x|z)$  is true posterior (decoder). The training objective of VAE is  $\mathbb{E}_{p^*(x)} [\mathcal{L}(x; \theta, \phi)]$  and it is optimized by SGVB with reparameterization trick. In vanilla VAE, prior  $p_\theta(z)$  is chosen as the standard Gaussian distribution.

Recently, some researchers show that the simplistic prior could lead to underfitting and many learnable priors are proposed subsequently to enrich prior [Tomczak and Welling, 2018]. Most of them focus on the aggregated posterior  $q_\phi(z)$ , which is shown as the optimal prior for ELBO by following decomposition where  $p_\lambda(z)$  denotes the prior maybe learnable:

$$\mathcal{L}(\theta, \phi, \lambda) = \mathbb{E}_{p^*(x)} \mathbb{E}_{q_\phi(z|x)} [\ln p_\theta(x|z)] + \mathbb{E}_{p^*(x)} [\mathbb{H}[q_\phi(z|x)]] + \mathbb{E}_{q_\phi(z)} \ln p_\lambda(z) = \mathcal{I} + \mathcal{J} + \mathcal{K} \quad (2)$$

where  $\mathcal{I}, \mathcal{J}, \mathcal{K}$  denotes 3 terms respectively. Notice that  $p_\lambda(z)$  only appears in the last term  $\mathcal{K}$  and the optimal solution of  $p_\lambda(z)$  is  $q_\phi(z)$ . However,  $q_\phi(z)$  is intractable. [Tomczak and Welling, 2018; Takahashi *et al.*, 2019] tries to approximate it but reaches limited performance.

## 2.2 GAN and Wasserstein distance

In vanilla GAN, a generator is trained to generate samples for deceiving discriminator, and a discriminator to distinguish the generated samples and real samples. However, vanilla GAN is unstable in training process. WGAN introduce Wasserstein distance to tackle this problem. 1st Wasserstein distance  $W^1(\mu, \nu)$  is used for calculating the distance between two measures  $\mu, \nu$ . The dual form of Wasserstein distance is following:

$$W^1(\mu, \nu) = \sup_{Lip(D) \leq 1} \{ \mathbb{E}_{\mu(x)} D(x) - \mathbb{E}_{\nu(x)} D(x) \} \quad (3)$$

where  $Lip(D) \leq 1$  means  $D$  is 1-Lipschitz. WGAN is optimized by minimizing  $W^1(p^*, p_\theta)$  which can be seen as a min-max optimization, whose parameters are  $D$  and  $\theta$ .

WGAN makes progress toward stable training but sometimes fails to converge due to the use of weight clipping to enforce the Lipschitz constrain. WGAN-GP [Gulrajani *et al.*, 2017] and WGAN-div [Wu *et al.*, 2018] improve WGAN by gradient penalty technique to achieve a more stable training. These techniques make WGAN framework become more robust and stable.

## 3 Pull-back Prior

### 3.1 Intuition of Pull-back Prior

The formula of Pull-back Prior is given by:

$$p_\lambda(z) = \frac{1}{Z} p_{\mathcal{N}}(z) \cdot e^{-\beta D(G(z))} \quad (4)$$

where  $p_{\mathcal{N}}$  is a simple prior,  $D$  is a discriminator,  $G$  is the generator  $G(z) = \mathbb{E}_{p_\theta(x|z)} x$ ,  $Z$  is the partition function  $Z = \int_{\mathcal{Z}} p_{\mathcal{N}}(z) \exp\{-\beta * D(G(z))\} dz$  and  $\beta$  is a learnable scalar.



Figure 1: The discriminators on above images (generated by linear interpolation of two sample from  $q_\phi(z)$ ), are better at both sides and worse at mid, which validates the intuition that discriminator can assess the quality of images. Moreover, the density of  $z$  which generate better images will increase, and the density of  $z$  which generate worse images will decrease.

An intuitive explanation of Pull-back Prior is given following: We would like to get a more powerful prior than simple prior  $p_{\mathcal{N}}$ . A simple way is to improve the density of  $z$  which generates better data and decrease the density of  $z$  which generates worse data.  $D$  is a discriminator to assess the quality of  $x$ . When  $D(x)$  is less,  $x$  is more similar to real data and of higher quality, as shown in fig. 1. We could pull-back the discriminator from data space to latent space, and function  $D(G(z))$  represents the quality of the data generated by  $z$ . To improve and decrease the density at better  $z$  and worse  $z$ , we modify  $p_{\mathcal{N}}(z)$  by  $\beta D(G(z))$  and then normalize it by  $Z$ , and finally we obtain the Pull-back Prior.

We obtain the basic formula of Pull-back Prior, and  $p_{\mathcal{N}}$  is a special case  $\beta = 0$ . The theoretical derivation for Pull-back Prior is provided in appendix A. However, it remains some troubles about how to obtain  $D$  and  $G$ , determine  $\beta$ , and calculate  $Z$ .

### 3.2 How to obtain $D$ and $G$

By definition in eq. (4),  $G$  is defined, and in practical networks, it is generated by neural network and  $p_\theta(x|z)$  set it as mean (e.g.,  $p_\theta(x|z)$  is Gaussian or Bernouli).

$D$  plays an important role in Pull-back Prior and We propose two way to obtain  $D$  and compare their difference in section 4.2 and section 4.1.

### 3.3 How to determine $\beta$

$\beta$  in eq. (4) represents how far  $p_\lambda$  is from  $p_{\mathcal{N}}$ , but how to decide the value of  $\beta$ ? When  $\beta$  is smaller, the difference between  $p_\lambda$  and  $p_{\mathcal{N}}$  is less, i.e., the influence of discriminator is severely limited. When  $\beta$  is larger,  $p_\lambda$  is farther from  $p_{\mathcal{N}}$ . Noticing that in eq. (11), we simplify the optimization of  $D$  by an approximated  $D$ , if  $p_\lambda$  is too far from  $p_{\mathcal{N}}$ , this approximation will become invalid. Consequently,  $\beta$  should be set to an appropriate value which can't severely limit the influence of discriminator and could ensure that approximated  $D$  is valid. It is important to realize that the Pull-back Prior is serving for better ELBO. Whatever the function family of  $p_\lambda$  is limited or approximation  $D$  is invalid, the ELBO will suffer. Therefore, it is reasonable to search  $\beta$  by the optimization for ELBO ( $\lambda$  contains  $\beta$  and  $\omega$ , which is the parameters of  $D$ ):

$$\beta = \arg \min_{\beta} \mathcal{L}(\theta, \phi, \lambda) = \arg \min_{\beta} \mathcal{L}(\theta, \phi, \beta, \omega) \quad (5)$$

The optimization process of  $\beta$  depends on  $\partial\mathcal{L}/\partial\beta$ :

$$\begin{aligned}\frac{\partial \ln Z}{\partial \beta} &= \frac{1}{Z} \int p_{\mathcal{N}}(z) e^{-\beta D(G(z))} \cdot (-D(G(z))) dz \\ &= \mathbb{E}_{p_{\lambda}(z)} [D(G(z))] \\ \frac{\partial \mathcal{L}}{\partial \beta} &= \mathbb{E}_{q_{\phi}(z)} [-D(G(z))] - \frac{\partial \ln Z}{\partial \beta} \\ &= -\mathbb{E}_{q_{\phi}(z)} [D(G(z))] + \mathbb{E}_{p_{\lambda}(z)} [D(G(z))] \quad (6)\end{aligned}$$

The 1st term in eq. (6) is the mean of discriminator on data generated from  $p_{\lambda}$ . The 2nd term in eq. (6) is the mean of discriminator on reconstructed data which is nearly same as real data when reconstruction is well-trained. Hence,  $\partial\mathcal{L}/\partial\beta = 0$  represents that the discriminator can't distinguish the reconstructed data (nearly same as real data) and generated data. It coincides the philosophy of GAN.

### 3.4 The lower bound of $Z$

The exact calculation of partition function  $Z$  is difficult. But in VAE domain, it is acceptable to obtain a lower-bound of log-likelihood. Therefore, we propose a feasible evaluation for upper-bound of  $Z$ , called  $\hat{Z}$ , which will not over-estimate log-likelihood. It is given by:

$$Z = \mathbb{E}_{q_{\phi}(z)} \frac{f_{\lambda}(z)}{q_{\phi}(z)} \leq \mathbb{E}_{q_{\phi}(z)} \frac{f_{\lambda}(z)}{\hat{q}_{\phi}(z)} = \hat{Z} \quad (8)$$

where  $f_{\lambda}(z) = p_{\mathcal{N}}(z) \exp\{-\beta D(G(z))\}$  and  $\hat{q}_{\phi}(z)$  is a lower-bound of  $q_{\phi}(z)$ . It is easy to show that

$$\begin{aligned}p_{\theta}(x) &= \int \frac{1}{Z} p_{\theta}(x|z) f_{\lambda}(z) dz \geq \int \frac{1}{\hat{Z}} p_{\theta}(x|z) f_{\lambda}(z) dz \\ \mathcal{K} &= \mathbb{E}_{q_{\phi}(z)} \frac{1}{Z} \ln f_{\lambda}(z) \geq \mathbb{E}_{q_{\phi}(z)} \frac{1}{\hat{Z}} \ln f_{\lambda}(z) = \hat{\mathcal{K}} \\ \mathcal{L} &= \mathcal{I} + \mathcal{J} + \mathcal{K} \geq \mathcal{I} + \mathcal{J} + \hat{\mathcal{K}} = \hat{\mathcal{L}}\end{aligned}$$

Above inequalities show that, if we replace  $Z$  by  $\hat{Z}$  we will obtain lower-bound of log-likelihood and ELBO. Such that  $\hat{Z}$  could be used in training and evaluation, and it will not over-estimate log-likelihood and ELBO.

The key of eq. (8) is the choice of  $\hat{q}_{\phi}(z)$ . As we mentioned before,  $q_{\phi}(z)$  is intractable to compute the exact density, but  $q_{\phi}(z|x)$  is feasible. We introduce a lower-bound of  $q_{\phi}(z)$ :

$$q_{\phi}(z) = \mathbb{E}_{p^*(x)} q_{\phi}(z|x) \approx \frac{1}{N} \sum_{i=1}^N q_{\phi}(z|x^{(i)}) \geq \frac{1}{N} q_{\phi}(z|x^{(j)})$$

where  $x^{(j)}$  is one of real data,  $N$  is the size of training set. To reduce the gap between  $\hat{q}_{\phi}(z)$  and  $q_{\phi}(z)$ ,  $q_{\phi}(z|x^{(j)})$  should be one of the largest in summation. Therefore, we firstly choose  $x^{(j)}$ , then sample  $z$  from  $q_{\phi}(z|x^{(j)})$ . By this way,  $q_{\phi}(z|x^{(j)})$  will be large enough.

In MNIST and other dataset consisting of Bernouli images,  $N$  might be very large. We introduce another ELBO which use  $q_{\phi}(z|e)$  instead of  $q_{\phi}(z|x)$ :

$$\begin{aligned}\mathbb{E}_{p^*(x)} \ln p_{\theta}(x) &\geq \mathbb{E}_{p^*(e)} \mathbb{E}_{p^*(x|e)} \ln \mathbb{E}_{q_{\phi}(z|e)} \frac{p_{\theta}(x|z) p_{\theta}(z)}{q_{\phi}(z|e)} \\ &= \mathbb{E}_{p^*(e)} \mathbb{E}_{p^*(x|e)} \mathbb{E}_{q_{\phi}(z|e)} \ln \frac{p_{\theta}(x|z) p_{\theta}(z)}{q_{\phi}(z|e)} \quad (7) \\ &= \mathbb{E}_{p^*(x)} \ln p^*(x) - \mathbb{E}_{p^*(e)} \mathbb{E}_{p^*(x|e)} KL(q_{\phi}(z|e), p_{\theta}(z|x))\end{aligned}$$

where  $p^*(x)$  is sampled from another distribution  $p^*(e)$ , and  $p^*(x|e)$  means the sampling process from  $e$ . The situation with finite  $N$  can be seen as a special case that  $p^*(x|e) = \delta(x - e)$ . eq. (7) is similar to the original ELBO eq. (1), and the above conclusion about learnable prior holds for eq. (7) by repeating above inference.

Because  $\beta$  is optimized from small to large during training, we use both estimations for  $Z$  in training. After training,  $\beta$  is large and  $p_{\lambda}$  approach to  $q_{\phi}$  by eq. (6), therefore we use eq. (8) for computing the final value of  $Z$ . The last thing we need to check is that the bias of estimation will not over-estimate the log-likelihood in evaluation:

$$p_{\theta}(x) = \mathbb{E}_{q_{\phi}(z|e)} \frac{f_{\lambda}(z) p_{\theta}(x|z)}{Z q_{\phi}(z|e)} \geq \mathbb{E}_{q_{\phi}(z|e)} \frac{f_{\lambda}(z) p_{\theta}(x|z)}{\hat{Z} q_{\phi}(z|e)}$$

which means the estimation log-likelihood based on  $\hat{Z}$  (RHS) is a lower bound of model density  $p_{\theta}(x)$ .

We have known that  $Z = \int_{\mathcal{Z}} p_{\mathcal{N}}(z) \exp\{-\beta * D(G(z))\} dz$ , denoted by  $\int_{\mathcal{Z}} f_{\lambda}(z) dz$ . It is natural to determine  $Z$  by importance sampling  $Z = \mathbb{E}_{p_{\mathcal{N}}(z)} \exp\{-\beta * D(G(z))\}$  as [Bauer and Mnih, 2019] did. By the theory of importance sampling, the variance of the estimation of  $Z$  is  $Var_{p_b}[\frac{f_{\lambda}}{p_b}]$ , where  $p_b$  is another distribution. Hence, the variance is smallest when  $p_b = p_{\lambda}$  and it is larger when  $p_{\lambda}$  is farther from  $p_b$ . If we choose  $p_b$  as  $p_{\mathcal{N}}$ , when  $\beta$  is large, the variance will be large and it will influence the optimization and evaluation.

The optimal choice for  $p_b$  is  $p_{\lambda}$  itself but it is hard to sample from  $p_{\lambda}$  during training. We try to find a distribution which is near to  $p_{\lambda}$  and easy-sampling. As eq. (6) shows, when  $\beta$  approaches optimal, the discriminator can't distinguish the data generated by  $p_{\lambda}(z)$  and  $q_{\phi}(z)$ . eq. (2) also shows that when  $p_{\lambda}(z)$  is optimized for  $\mathcal{L}(\theta, \phi, \lambda)$ , it approaches to  $q_{\phi}$ . Consequently, it is reasonable to choose  $q_{\phi}$  as  $p_b$ .

Since  $q_{\phi}(z|e)$  is known, above bias estimation of  $q_{\phi}(z)$  is feasible by  $\frac{1}{N} q_{\phi}(z|e^{(j)})$ .  $\hat{q}_{\phi}(z)$  denotes the bias estimation of  $q_{\phi}(z)$ . Then, a bias estimation  $\hat{Z}$  is given by:

$$Z = \mathbb{E}_{q_{\phi}(z)} \frac{f_{\lambda}(z)}{q_{\phi}(z)} \leq \mathbb{E}_{q_{\phi}(z)} \frac{f_{\lambda}(z)}{\hat{q}_{\phi}(z)} = \hat{Z} \quad (8)$$

Because  $\beta$  is optimized from small to large during training, we use both estimations for  $Z$  in training. After training,  $\beta$  is large and  $p_{\lambda}$  approach to  $q_{\phi}$  by eq. (6), therefore we use eq. (8) for computing the final value of  $Z$ . The last thing we need to check is that the bias of estimation will not over-estimate the log-likelihood in evaluation:

$$p_{\theta}(x) = \mathbb{E}_{q_{\phi}(z|e)} \frac{f_{\lambda}(z) p_{\theta}(x|z)}{Z q_{\phi}(z|e)} \geq \mathbb{E}_{q_{\phi}(z|e)} \frac{f_{\lambda}(z) p_{\theta}(x|z)}{\hat{Z} q_{\phi}(z|e)}$$

which means the estimation log-likelihood based on  $\hat{Z}$  (RHS) is a lower bound of model density  $p_{\theta}(x)$ .

## 4 Training and Sampling

In this section, we will propose two training methods and a sampling method for VAEPP. The main difference between these two training method is the selection of discriminator.

---

**Algorithm 1** Naive training algorithm for VAEPP

---

**Require:** The gradient penalty algorithm  $R$ , the batch size  $b$ , the number of critic iterations per generator iteration  $n_c$ , the parameters for Adam Optimizers,  $\tau$ .

```

1: while  $\theta, \phi, \beta, \omega$  have not converged do
2:   for  $k = 1, \dots, n_c$  do
3:     for  $i = 1, \dots, b$  do
4:       Sample  $e, x \sim p^*, z \sim q_\phi(z|e), \epsilon \sim p_N$ 
5:        $Z^{(i)} \leftarrow \frac{1}{2}(\exp\{-\beta D(G(\epsilon))\} + \frac{f_\lambda(z)}{q_\phi(z)})$ 
6:        $\mathcal{L}^{(i)} \leftarrow \ln p_\theta(x|z) + \ln f_\lambda(z) - \ln q_\phi(z|e)$ 
7:     end for
8:      $\mathcal{L} \leftarrow \frac{1}{b} \sum_i \mathcal{L}^{(i)} - \ln(\frac{1}{b} \sum_i Z^{(i)})$ 
9:      $\theta, \phi, \beta \leftarrow \text{Adam}(\nabla_{\theta, \phi, \beta} \mathcal{L}, \{\theta, \phi, \beta\}, \tau)$ 
10:  end for
11:  for  $i = 1, \dots, b$  do
12:    Sample  $e, x \sim p^*$ , latent variable  $z \sim p_N$ 
13:     $\hat{x} = \mathbb{E}_{p_\theta(x|z)}[x]$ , get gradient penalty  $\zeta \leftarrow R(e, \hat{x})$ 
14:     $L^{(i)} \leftarrow D(\hat{x}) - D(x) + \zeta$ 
15:  end for
16:   $\omega \leftarrow \text{Adam}(\nabla_\omega \frac{1}{b} \sum_i L^{(i)}, \omega, \tau)$ 
17: end while

```

---




Figure 2: Training process of Naive VAEPP and VAEPP on CIFAR-10. Naive VAEPP is more unstable and nearly crash at 80 epoch while VAEPP has a little acceptable gap. From global view, the training loss of VAEPP is more smooth than Naive VAEPP and is better than Naive VAEPP over almost all training process, which validates the motivation in section 4.2. There are little gaps at per 200 epoch because learning rate is reduced to half at every 200 epoch.

#### 4.1 Naive training for VAEPP

Discriminator is selected by  $W^1(p^\dagger, p^*)$  where  $p^\dagger(x) = \mathbb{E}_{p_N(z)} p_\theta(x|z)$ , and it could be trained directly by WGAN-GP or WGAN-div. Naive training algorithm for VAEPP is provided in algorithm 1, where the training of WGAN and SGVB runs alternately.

#### 4.2 Combing training for VAEPP

However, we notice that in the training process, the optimization of  $\omega$  may influence the optimization of  $\theta, \phi, \beta$ , *e.g.*, the optimization for  $\omega$  significantly worsen the loss function, as show in fig. 2. The reason is that the optimization for  $\omega$  is independent to the optimization of  $\theta, \phi, \beta$  in algorithm 1. This independence is from the philosophy of GAN but may lower the performance of VAEPP (log-likelihood). Hence, it is necessary to combine these two optimization into one to improve the performance and stability of VAEPP. Our solution is to use SGVB with gradient penalty regularizer to train VAEPP, *i.e.*,  $\max_{\theta, \phi, \beta} \max_{Lip(D) \leq 1} \mathcal{L}(\theta, \phi, \beta, \omega)$ .

---

**Algorithm 2** Combing training algorithm for VAEPP

---

**Require:** The gradient penalty algorithm  $R$ , the batch size  $b$ , the parameters for Adam Optimizers,  $\tau$ .

```

1: while  $\theta, \phi, \beta, \omega$  have not converged do
2:   for  $i = 1, \dots, b$  do
3:     Sample  $e, x \sim p^*, z \sim q_\phi(z|e), \epsilon \sim p_N$ 
4:      $\hat{x} = \mathbb{E}_{p_\theta(x|e)}[x]$ , get gradient penalty  $\zeta \leftarrow R(e, \hat{x})$ 
5:      $Z^{(i)} \leftarrow \frac{1}{2}(\exp\{-\beta D(G(\epsilon))\} + \frac{f_\lambda(z)}{q_\phi(z)})$ 
6:      $\mathcal{L}^{(i)} \leftarrow \ln p_\theta(x|z) + \ln f_\lambda(z) - \ln q_\phi(z|e) + \beta \zeta$ 
7:   end for
8:    $\mathcal{L} \leftarrow \frac{1}{b} \sum_i \mathcal{L}^{(i)} - \ln(\frac{1}{b} \sum_i Z^{(i)})$ 
9:    $\theta, \phi, \beta, \omega \leftarrow \text{Adam}(\nabla_{\theta, \phi, \beta} \mathcal{L}, \{\theta, \phi, \beta, \omega\}, \tau)$ 
10: end while

```

---

In such optimization, the behavior of  $\theta, \phi, \beta$  is same as algorithm 1 since the optimization for them is same. For analysis of  $\omega$ , we firstly show an inequality of  $\ln Z$ :

$$\ln Z = \ln \mathbb{E}_{p_N(z)} e^{-\beta * D(G(z))} \geq \mathbb{E}_{p_N(z)} [-\beta * D(G(z))]$$

Then  $\max_{Lip(D) \leq 1} \mathcal{L}(\theta, \phi, \beta, \omega)$  indeed find a suboptimal solution for  $W^1(p^\dagger, p^*)$  (sign  $\simeq$  means that optimizations at left and right are equivalent):

$$\begin{aligned}
\max_{Lip(D) \leq 1} \mathcal{L} &\simeq \max_{Lip(D) \leq 1} \{-\mathbb{E}_{q_\phi(z)} \beta * D(G(z)) - \ln Z\} \\
&\leq \beta \max_{Lip(D) \leq 1} \{\mathbb{E}_{p_N(z)} D(G(z)) - \mathbb{E}_{q_\phi(z)} D(G(z))\} \quad (9) \\
&= \beta W^1(p^\dagger, p_r) \approx \beta W^1(p^\dagger, p^*)
\end{aligned}$$

where  $p_r$  denotes  $p_r(x) = \mathbb{E}_{q_\phi(z)} p_\theta(x|z)$ , consisting of reconstructed data. The last approximation sign is from the fact that  $p_r \rightarrow p^*$  after a few epoch in the training of VAE. eq. (9) also uses the assumption  $D(G(z)) = \mathbb{E}_{p_\theta(x|z)} D(x)$  introduced in appendix A to simplify equation.

eq. (9) indicates that it is reasonable to gain a suboptimal solution for  $D$  by directly optimizing  $\mathcal{L}$ , and the gradient penalty term should be multiplied by  $\beta$ . By this way, the optimizations for  $\omega$  and  $\theta, \phi, \beta$  is combined into one, which provides as algorithm 2. Thanks to the stable and efficient gradient penalty regularizer term provided by WGAN-GP and WGAN-div, we enjoy stable and efficient training. We will call the model trained by algorithm 1 Naive VAEPP and the model trained by algorithm 2 VAEPP.

#### 4.3 Sampling from VAEPP

It is not easy to sample  $z$  from  $p_\lambda(z)$  since the formula of  $p_\lambda(z)$  is complicated. Accept/Reject Sampling (ARS) is also not useful for  $p_\lambda$  because ARS requires that  $p_\lambda(z)/p_N(z)$  is bounded by a constant  $M$  (It means  $\beta$  is limited to a very small value), such that a sample could be sampled in  $M$  times.

Langevin Dynamics may be a useful sampling method because it only requires that  $\nabla_z \log p_\lambda(z)$  is computable and the initial  $z_0$  has an enough high density [Song and Ermon, 2019]. Moreover, MEG [Kumar *et al.*, 2019] have implemented a Metropolis-Adjusted Langevin Algorithm (MALA) for sampling where the formula of density is similar to  $p_\lambda$  and also contains a discriminator term. But the selection of initial  $z_0$  whose density is high enough is still a problem.

Model	MNIST	CIFAR
<b>With autoregressive</b>		
PixelCNN	81.30	3.14
DRAW	80.97	3.58
IAFVAE	79.88	3.11
PixelVAE++	78.00	2.90
PixelRNN	79.20	3.00
VLA	79.03	2.95
PixelSNAIL		2.85
PixelHVAE with VampPrior	78.45	
<b>Without autoregressive</b>		
Implicit Optimal Priors	83.21	
Discrete VAE	81.01	
LARS	80.30	
VampPrior	79.75	
BIVA	78.59	3.08
<b>Naive VAEPP</b>	76.49	3.15
<b>VAEPP</b>	76.37	2.91
<b>VAEPP+Flow</b>	76.23	2.84

Table 1: Test NLL on MNIST and Bits/dim on CIFAR-10. Bits/dim means  $-\log p_\theta(x|z)/(3072 * \ln(2))$ . The data is from [Maaløe *et al.*, 2019], [Chen *et al.*, 2018], [Tomczak and Welling, 2018], [Bauer and Mnih, 2019] and [Takahashi *et al.*, 2019]. VAEPP+Flow means VAEPP with a normalization flow on encoder, to enhance the ability of encoder. Additional, we compare VAE based on  $q_\phi(z|x)$  and  $q_\phi(z|e)$  on MNIST, whose NLL are 81.10 and 83.30 respectively. Moreover, evaluation using importance sampling based on  $q_\phi(z|e)$  has small enough variance (0.01) with  $10^3$  samples from  $q_\phi(z|e)$  and 10 samples for  $p^*(x|e)$ . It validates that using  $q_\phi(z|e)$  is stable for evaluation and don’t improve the performance. VAEPP reaches the state of art, and is competitive to the models with autoregressive component.

Following the philosophy of VAEPP, *i.e.*, using the technique of GAN to assist VAE, it is natural to use GAN to model the distribution  $q_\phi(z)$ , and use samples of GAN as the initial point of MALA, which has high enough density in  $p_\lambda(z)$ . The sampling of VAEPP consists of 3 parts: generate initial  $z_0$  by a GAN, then generate  $z \sim p_\lambda(z)$  by Langevin Dynamics, and finally generate  $x$  by decoder. This sampling process is similar to 2-Stage VAE [Dai and Wipf, 2019]. The main difference between them is that the prior of VAEPP is explicit, but 2-Stage VAE not. Therefore, Langevin Dynamics is applied in VAEPP to sample from the explicit prior. In experiments, sampling from the explicit learnable prior improves the quality of sampling and ensures theoretical correctness of prior.

## 5 Experiments

VAEPP is evaluated in vast common datasets including MNIST, Fashion-MNIST [Xiao *et al.*, 2017], Omniglot [Lake *et al.*, 2015], CIFAR-10 [Krizhevsky *et al.*, 2009] and CelebA [Liu *et al.*, 2015] with metrics log-likelihood, FID [Heusel *et al.*, 2017] and IS [Salimans *et al.*, 2016] to show the performance of VAEPP.

Model	Static MNIST	Fashion	Omniglot
Naive VAEPP	78.06	214.63	90.72
VAEPP	77.73	213.24	89.60
VAEPP+Flow	77.66	213.19	89.24

Table 2: Test NLL on Static MNIST, Fashion-MNIST and Omniglot.

GP Strategy	Naive VAEPP	VAEPP
WGAN-GP	3.15	2.95
WGAN-div-1	3.20	2.91
WGAN-div-2	4.47	2.99

Table 3: Comparison between VAEPP and Improved VAEPP when gradient penalty strategy varies on CIFAR-10 with  $\dim \mathcal{Z} = 1024$ . For any gradient penalty strategy in the table, VAEPP outperforms Naive VAEPP, which validates the our intuition of design of algorithm 2. We select WGAN-div-1 as our default gradient penalty strategy since it reaches best performance in VAEPP.

### 5.1 Log-likelihood Evaluatoin

We evaluate and compare the performance of VAEPP trained by algorithm 1 and algorithm 2 on CIFAR10, as the gradient penalty algorithm is selected from 3 strategies: WGAN-GP, WGAN-div-1 (sampling the linear combination of two real or two fake data points) and WGAN-div-2 (sampling both real or fake data points), as shown in table 3. Our conclusion is that algorithm 2 outperforms algorithm 1 under all of settings in CIFAR-10 and we select WGAN-div-1 as default setting.

We compare our algorithms with other log-likelihood based model on MNIST and CIFAR-10 as shown in table 1, and another table 2. Because the improvement of autoregressive components is significant, we separate models by whether use auto-regressive component as [Maaløe *et al.*, 2019] did. VAEPP outperforms most of the models without autoregressive component and is competitive to the models with autoregressive component. The reason of why VAEPP don’t use auto-regressive component is that VAEPP is time-consuming in training, evaluation and sampling due to the huge structure (need additional discriminator) and Langevin Dynamics. It is not easy to apply auto-regressive component on VAEPP considering auto-regressive component is also time-consuming. Therefore, how to apply autoregressive component on VAEPP is a valuable and challenging practical work. We leave it as a future work.

To validate that it is better to use  $q_\phi(z)$  to evaluate  $Z$  than  $p_{\mathcal{N}}(z)$  in section 3.4, we calculate the  $KL(q_\phi(z)||p_\lambda(z))$  and  $KL(p_{\mathcal{N}}(z)||p_\lambda(z))$  on CIFAR-10 and MNIST. The former is less than the difference (180.3 on CIFAR-10 and 12.497 on MNIST) between reconstruction term and ELBO [Hoffman and Johnson, 2016], and the latter one can be evaluated directly (1011.30 on CIFAR-10 and 57.45 on MNIST). Consequently,  $q_\phi(z)$  is much closer to  $p_\lambda(z)$  than  $p_{\mathcal{N}}(z)$ .

To ensure the variance of estimation  $\hat{Z}$  is small enough, the  $q_\phi(z|e)$  is selected as truncated normal distribution (drop the sample whose magnitude is more than 2 standard deviation from the mean) instead of normal distribution. In eq. (8),



Figure 3: Examples of generated images.

Model	MNIST	Fashion	CIFAR	CelebA
Best GAN	$\sim 10$	$\sim 32$	$\sim 70$	$\sim 49$
VAE+Flow	54.8	62.1	81.2	65.7
WAE-MMD	115.0	101.7	80.9	62.9
2-StageVAE	12.6	29.3	72.9	44.4
GAN-VAEPP	12.7	28.1	74.1	53.4
VAEPP	12.0	28.2	71.0	53.4

Table 4: FID comparisons to GAN-based models and other VAEs. Best GAN indicates the best FID on each dataset across all GAN models when trained using settings suggested by original authors. VAEPP uses Bernoulli as posterior on MNIST and Discretized Logistic [Salimans *et al.*, 2017] on others. GAN-VAEPP indicates that image is directly sampled from  $z_0$ , generated by GAN without Langevin Dynamics. The data of Best GAN and other VAEs is from [Dai and Wipf, 2019]. The FID of VAEPP is usually better than GAN in VAEPP (except Fashion, the texture is hard to reconstruct, which leads to the reconstruction is much different to real image. It will misguide the Langevin Dynamics since we have assumed that), which validates that the explicit prior and Langevin Dynamics are useful for improving the quality of sampling.

$\hat{q}_\phi(z)$  estimated by  $\frac{1}{N}q_\phi(z|e^{(j)})$  is the denominator and if it becomes too small sometimes, the estimation will become unstable. If  $q_\phi(z|e^{(j)})$  is selected as normal distribution, the probability of  $z$  in tail will be larger when the number of samples become larger, which will lead to huge variance to the estimation  $\hat{Z}$ . With  $10^9$  samples, the variance of these two method are 0.809260 (normal) and 0.000967 (truncated normal) in MNIST. Therefore, truncated normal is selected as default setting.

## 5.2 Quality of Sampling

As common sense, the quality of samples of VAE is worse than GAN, and it is indeed a reason that we involve the techniques of GAN to improve the VAE model: We use discriminator to adjust learnable prior and GAN to sample the initial  $z_0$  for Langevin Dynamics. These techniques will help VAEPP improve the quality of samples. The samples of VAEPP gets good FID and IS, competitive to GANs and 2-Stage VAE (which is the SOTA of VAE in FID), as shown in table 4. Some generated images is shown in fig. 3.

It is hard to reach best performance in FID, IS and log-likelihood simultaneously by same setting. We observe this fact that when dimension of latent space is increasing, the trend of FID, IS is different to the trend of log-likelihood, as shown in fig. 4. As diagnosis in [Dai and Wipf, 2019], the variance of  $p_\theta(x|z)$  is chosen as a learnable scalar, and the dimension of latent space is selected as a number that little larger than the dimension of real data manifold,  $\dim \mathcal{Z} = 128$ , as our experimental result.




Figure 4: Comparison of VAEPP with learnable scalar variance of  $p_\theta(x|z)$ , as the dimension of latent space varies on CIFAR-10, with metrics BPD, FID and IS. FID and BPD is better when it is smaller and IS is better when it is larger. When dimension of latent space is greater than 128, the quality of sampling becomes worse and BPD becomes better as  $\dim \mathcal{Z}$  increases. It validates the proposition that  $\dim \mathcal{Z}$  should be selected as a minimal number of active latent dimensions in [Dai and Wipf, 2019]. The reconstruction term is optimized more as  $\dim \mathcal{Z}$  is larger, because larger latent space could remain more information. Meanwhile, the  $KL(q_\phi(z)||p_\lambda(z))$  (Bounded by the difference between ELBO and reconstruction term) will not increase very much since learnable prior minimizes it. It also shows a phenomenon that FID, IS is not always keep same trend as BPD, maybe greatly different.

For better understanding, the value of discriminator in this section is normalized into  $\mathcal{N}(0, 1)$  in training set.

To validate the eq. (6), we calculate the  $\mathbb{E}_{p_\lambda(z)}[D(G(z))]$  (discriminator on generated samples) and  $\mathbb{E}_{q_\phi(z)}[D(G(z))]$  (discriminator on reconstructed samples). They are 0.092 and 0.015 respectively on CIFAR-10, which means discriminator on generated samples and reconstructed samples are nearly same as discriminator on real data.

To validate the assumption in appendix A holds in actual experiment, we calculate  $|\mathbb{E}_{p_\theta(x|z)} D(x) - D(G(z))|$  (0.019) on CIFAR-10, which is acceptable small.

## 6 Conclusion

We propose a novel learnable prior Pull-back Prior for VAE model, by adjusting prior with a discriminator assessing the quality of data, with solid inference and intuitive explanation. We propose an efficient and stable training method for VAE with Pull-back Prior, by mixing the optimizations of WGAN and VAE into one. VAEPP is evaluated on vast common dataset, and shows impressive performance in log-likelihood and quality of sampling. We believe that this paper could lead VAE models into a new stage, with clear formula, general framework and powerful performance.

## A Derivation of Pull-back Prior

Reviewing the inference of aggregated posterior, the optimization  $\min_{\theta, \phi, \lambda} \mathcal{L}(\theta, \phi, \lambda)$  could be deformed into  $\min_{\theta, \phi} J(\theta, \phi)$ , where  $J(\theta, \phi) = \min_{\lambda} \mathcal{L}(\theta, \phi, \lambda)$ .  $J(\theta, \phi)$  means to search the analytical optimal prior  $\lambda$  when  $\theta, \phi$

is given. When  $\mathcal{L}$  is ELBO, the optimal solution in  $J(\theta, \phi)$  is  $p_\lambda(z) = q_\phi(z)$ . Our idea is to set another objective function  $\hat{\mathcal{L}}$  for  $J(\theta, \phi)$ , *i.e.*, search another optimal prior  $\lambda$  for  $J(\theta, \phi) = \min_\lambda \hat{\mathcal{L}}(\theta, \phi, \lambda)$ . Because ELBO is derived by  $KL(p_\theta, p^*)$ , a candidate  $\hat{\mathcal{L}}$  could be Wasserstein distance,  $W^1(p_\theta, p^*)$ . Considering following optimization:

$$J(\theta, \phi) = \min_\lambda \hat{\mathcal{L}}(\theta, \phi, \lambda) = \min_\lambda W^1(p_\theta, p^*) = \min_\lambda \sup_{\text{Lip}(D) \leq 1} \{ \mathbb{E}_{p_\lambda(z)} \mathbb{E}_{p_\theta(x|z)} D(x) - \mathbb{E}_{p^*(x)} D(x) \} \quad (10)$$

It is hard to get an analytical solution of  $\lambda$  directly from eq. (10), therefore we add an assumptions and an approximation to simplify it. Since  $p_\theta(x|z)$  is usually a distribution with small variance, it is rational to assume  $\mathbb{E}_{p_\theta(x|z)} D(x) = D(G(z))$  (It is an extension definition when posterior is Bernouli). Even though, the optimization  $\sup_{\text{Lip}(D) \leq 1}$  is still tough because the optimization of  $D$  is independent on  $\lambda$ . If we restrict  $p_\lambda$  near the  $p_{\mathcal{N}}$ , an approximation  $D$  could be used to replace it, *e.g.*, in section 4.1 or section 4.2. Consequently, the simplified optimization is following:

$$\min_\lambda \{ \mathbb{E}_{p_\lambda(z)} D(G(z)) - \mathbb{E}_{p^*(x)} D(x) \} \quad (11)$$

**s.t.**  $KL(p_\lambda, p_{\mathcal{N}}) = \alpha, \quad \int_{\mathcal{Z}} p_\lambda(z) dz = 1$

It could be solved by Lagrange multiplier method introduced by calculus of variation [Gelfand *et al.*, 2000]. The Lagrange function with Lagrange multiplier  $\eta, \gamma$  is following:

$$F(p_\lambda, \eta, \gamma) = \mathbb{E}_{p_\lambda(z)} D(G(z)) - \mathbb{E}_{p^*(x)} D(x) + \eta \left( \int_{\mathcal{Z}} p_\lambda(z) dz - 1 \right) + \gamma (KL(p_\lambda, p_{\mathcal{N}}) - \alpha) \quad (12)$$

We solve eq. (12) by Euler-Lagrange equation:

$$D(G(z)) + \eta + \gamma(\ln p_\lambda(z) + 1 - \ln p_{\mathcal{N}}(z)) = 0 \quad (13)$$

Therefore,  $\ln p_\lambda(z) = \frac{1}{\gamma} D(G(z)) + \ln p_{\mathcal{N}}(z) + (\frac{\eta}{\gamma} - 1)$  is the optimal solution, where  $\gamma$  is determined by  $\alpha$  and  $\eta$  is determined from condition  $\int_{\mathcal{Z}} p_\lambda(z) dz = 1$ .

Consequently,  $\beta$  is determined by  $\alpha$ , representing how far  $p_\lambda$  is from  $p_{\mathcal{N}}$ . In eq. (11),  $\alpha$  is static and should be selected as an appropriate value, *i.e.*,  $\beta$  should be searched, as section 3.3 does. eq. (13) holds for any discriminator  $D$ , and therefore the selection of  $D$  plays an important role, which is discussed in section 4.1 and section 4.2.

## References

- [Bauer and Mnih, 2019] Matthias Bauer and Andriy Mnih. Resampled priors for variational autoencoders. In *The 22nd International Conference on Artificial Intelligence and Statistics*, pages 66–75, 2019.
- [Chen *et al.*, 2018] Xi Chen, Nikhil Mishra, Mostafa Rohaninejad, and Pieter Abbeel. Pixelsnail: An improved autoregressive generative model. In *International Conference on Machine Learning*, pages 863–871, 2018.
- [Dai and Wipf, 2019] Bin Dai and David Wipf. Diagnosing and enhancing vae models. *arXiv preprint arXiv:1903.05789*, 2019.
- [Dinh *et al.*, 2016] Laurent Dinh, Jascha Sohl-Dickstein, and Samy Bengio. Density estimation using real nvp. *arXiv preprint arXiv:1605.08803*, 2016.
- [Gelfand *et al.*, 2000] Izrail Moiseevitch Gelfand, Richard A Silverman, et al. *Calculus of variations*. Courier Corporation, 2000.
- [Goodfellow *et al.*, 2014] Ian Goodfellow, Jean Pouget-Abadie, Mehdi Mirza, Bing Xu, David Warde-Farley, Sherjil Ozair, Aaron Courville, and Yoshua Bengio. Generative adversarial nets. In *Advances in neural information processing systems*, pages 2672–2680, 2014.
- [Gulrajani *et al.*, 2017] Ishaan Gulrajani, Faruk Ahmed, Martin Arjovsky, Vincent Dumoulin, and Aaron C Courville. Improved training of wasserstein gans. In *NIPS*, 2017.
- [Heusel *et al.*, 2017] Martin Heusel, Hubert Ramsauer, et al. Gans trained by a two time-scale update rule converge to a local nash equilibrium. In *Advances in Neural Information Processing Systems*, pages 6626–6637, 2017.
- [Hoffman and Johnson, 2016] Matthew D Hoffman and Matthew J Johnson. Elbo surgery: yet another way to carve up the variational evidence lower bound. In *Workshop in Advances in Approximate Bayesian Inference, NIPS*, volume 1, 2016.
- [Kingma and Dhariwal, 2018] Durk P Kingma and Prafulla Dhariwal. Glow: Generative flow with invertible 1x1 convolutions. In *Advances in Neural Information Processing Systems*, pages 10215–10224, 2018.
- [Kingma and Welling, 2014] Diederik P Kingma and Max Welling. Auto-encoding variational bayes. In *ICLR*, 2014.
- [Kingma *et al.*, 2016] Durk P Kingma, Tim Salimans, Rafal Jozefowicz, Xi Chen, Ilya Sutskever, and Max Welling. Improved variational inference with inverse autoregressive flow. In *Advances in neural information processing systems*, pages 4743–4751, 2016.
- [Krizhevsky *et al.*, 2009] Alex Krizhevsky, Geoffrey Hinton, et al. Learning multiple layers of features from tiny images. Technical report, Citeseer, 2009.
- [Kumar *et al.*, 2019] Rithesh Kumar, Anirudh Goyal, Aaron Courville, and Yoshua Bengio. Maximum entropy generators for energy-based models. *arXiv preprint arXiv:1901.08508*, 2019.



- [Lake *et al.*, 2015] Brenden M Lake, Ruslan Salakhutdinov, and Joshua B Tenenbaum. Human-level concept learning through probabilistic program induction. *Science*, 350(6266):1332–1338, 2015.
- [Liu *et al.*, 2015] Ziwei Liu, Ping Luo, Xiaogang Wang, and Xiaoou Tang. Deep learning face attributes in the wild. In *Proceedings of the IEEE international conference on computer vision*, pages 3730–3738, 2015.
- [Maaløe *et al.*, 2019] Lars Maaløe, Marco Fraccaro, Valentin Liévin, and Ole Winther. Biva: A very deep hierarchy of latent variables for generative modeling. *arXiv preprint arXiv:1902.02102*, 2019.
- [Rezende *et al.*, 2014] Danilo Jimenez Rezende, Shakir Mohamed, and Daan Wierstra. Stochastic backpropagation and approximate inference in deep generative models. In *ICML*, 2014.
- [Salimans *et al.*, 2016] Tim Salimans, Ian Goodfellow, Wojciech Zaremba, Vicki Cheung, Alec Radford, and Xi Chen. Improved techniques for training gans. In *Advances in neural information processing systems*, pages 2234–2242, 2016.
- [Salimans *et al.*, 2017] Tim Salimans, Andrej Karpathy, Xi Chen, and Diederik P Kingma. Pixelcnn++: Improving the pixelcnn with discretized logistic mixture likelihood and other modifications. *arXiv preprint arXiv:1701.05517*, 2017.
- [Song and Ermon, 2019] Yang Song and Stefano Ermon. Generative modeling by estimating gradients of the data distribution. In *Advances in Neural Information Processing Systems*, pages 11895–11907, 2019.
- [Takahashi *et al.*, 2019] Hiroshi Takahashi, Tomoharu Iwata, Yuki Yamanaka, Masanori Yamada, and Satoshi Yagi. Variational autoencoder with implicit optimal priors. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 33, pages 5066–5073, 2019.
- [Tomczak and Welling, 2016] Jakub M Tomczak and Max Welling. Improving variational auto-encoders using householder flow. *arXiv preprint arXiv:1611.09630*, 2016.
- [Tomczak and Welling, 2018] Jakub Tomczak and Max Welling. Vae with a vampprior. In *International Conference on Artificial Intelligence and Statistics*, pages 1214–1223, 2018.
- [Van den Oord *et al.*, 2016] Aaron Van den Oord, Nal Kalchbrenner, Lasse Espeholt, Oriol Vinyals, Alex Graves, et al. Conditional image generation with pixelcnn decoders. In *Advances in neural information processing systems*, pages 4790–4798, 2016.
- [Wu *et al.*, 2018] Jiqing Wu, Zhiwu Huang, Janine Thoma, Dinesh Acharya, and Luc Van Gool. Wasserstein divergence for gans. In *Proceedings of the European Conference on Computer Vision (ECCV)*, pages 653–668, 2018.
- [Xiao *et al.*, 2017] Han Xiao, Kashif Rasul, and Roland Vollgraf. Fashion-mnist: a novel image dataset for benchmarking machine learning algorithms, 2017.